

Recommender Systems Workshop

Slideshow: <https://bit.ly/3ji05zy>

Notebook: <https://bit.ly/3dzMDmM>

Presented by Jack Douglas

Introduction



Jack Douglas

(he/him)

2B Software Engineering

Email: jack.douglas@uwaterloo.ca

LinkedIn: <https://bit.ly/3gWuXUo>

Workshop Overview

Intended Audience: Beginner to Intermediate Data Scientists

Goals:


- Learn about what recommender systems do and the types of data that you will see in the real world
- Learn about the two main types of recommender systems
- Learn how to implement a recommender system on your own

Workshop Outline

1. What is a Recommender System?
2. Motivation
3. How Recommender Systems Work
4. Types of Data
5. Useful Notation
6. Approaches to Recommendations
 - a. Content-based Filtering
 - b. Collaborative Filtering
7. Movie Recommendations Notebook
8. Deep Learning Recommender Systems


What is a Recommender System?

- **Def'n:** a type of machine learning system that recommend new products and services to users
- Recommender system looks at patterns of activities between different users and different products to produce these recommendations



Motivation

- Recommendation systems are used everywhere
 - Advertisements/product recommendations, song recommendations, and movie recommendations are some examples
- Understanding recommender systems is a practical skill that can help make your résumé stand out!



How Recommender Systems Work

- Recommender systems seek to understand 3 key relationships:
 - **User - Product**: based on users' individual product preferences
 - **Product - Product**: based on similar products, either by appearance or description
 - **User - User**: based on similar users' likely having similar product preferences
- A combination of these relationships is used to sort the best recommendations for users

User - Product Example

If a user rates a movie highly, then the recommender system uses the connection between the user and genre, cast, director, etc. to create recommendations:



User 1



User 1 Profile

- Likes comedy
- Likes Will Ferrell
- Likes John C. Reilly
- etc.




Product - Product Example

If two products share many attributes, then the recommender system forms a connection between the two products:



Similar casts, similar genres, same director




User - User Example

If a two users rates a movie highly, then the recommender system uses that connection between the two to help create recommendation:



User 1



User 2



Type of Data: User Behaviour Data

- **Explicit Ratings:** ratings provided by the user
 - Ex. ratings, reviews, likes, following
- **Implicit Ratings:** ratings provided from users interacting with a product
 - Ex. clicks, views, purchases
- Used in every type of recommender system

Movie	User 1	User 2
Step Brothers (1)	5	1
Anchorman (2)	?	1
The Other Guys (3)	?	1
Taxi Driver (4)	4	?
Goodfellas (5)	5	5

Explicit Ratings Example

Type of Data: User Demographic Data

- Examples: age, location, education, income
- Used in some recommender systems, can help improve the recommendations



NA Adapter



UK Adapter

Type of Data: Product Attributes

- Examples: genre, type of item, company producing item
- Needed in content-based filtering



Attributes: Nike, T-Shirt, Black,
Graphic, etc.

Movie Rating Dataset Examples

—

Movie	User 1	User 2	x_1 (comedy)	x_2 (drama)
Step Brothers (1)	5	1	1	0
Anchorman (2)	?	1	1	0
The Other Guys (3)	?	1	1	0
Taxi Driver (4)	4	?	0	1
Goodfellas (5)	5	5	0	1

Dataset with explicit ratings and product attributes

Useful Notation

$r(i, j)$: 1 if user j has rated movie i (0 otherwise)

$x^{(i)}$: feature vector for movie i (first index is always 1)

$\theta^{(j)}$: parameter vector for user j

$y^{(i,j)}$: actual rating by user j on movie i (if defined)

$(\theta^{(j)})^T x^{(i)}$: predicted rating by user j on movie i

Notation Examples

Movie	User 1	User 2	x_1 (comedy)	x_2 (drama)
Step Brothers (1)	5	1	1	0
Anchorman (2)	?	1	1	0
The Other Guys (3)	?	1	1	0
Taxi Driver (4)	4	?	0	1
Goodfellas (5)	5	5	0	1

$r(i, j)$: 1 if user j has rated movie i (0 otherwise)

$x^{(i)}$: feature vector for movie i (first index is always 1)

$\theta^{(j)}$: parameter vector for user j

$y^{(i,j)}$: actual rating by user j on movie i (if defined)

$(\theta^{(j)})^T x^{(i)}$: predicted rating by user j on movie i

$$r(1,1) = 1$$

$$r(4,2) = 0$$

$$x^{(2)} = \begin{bmatrix} 1 \\ 1 \\ 0 \end{bmatrix}$$

$$\theta^{(1)} = \begin{bmatrix} 5 \\ 0 \\ 5 \end{bmatrix}$$

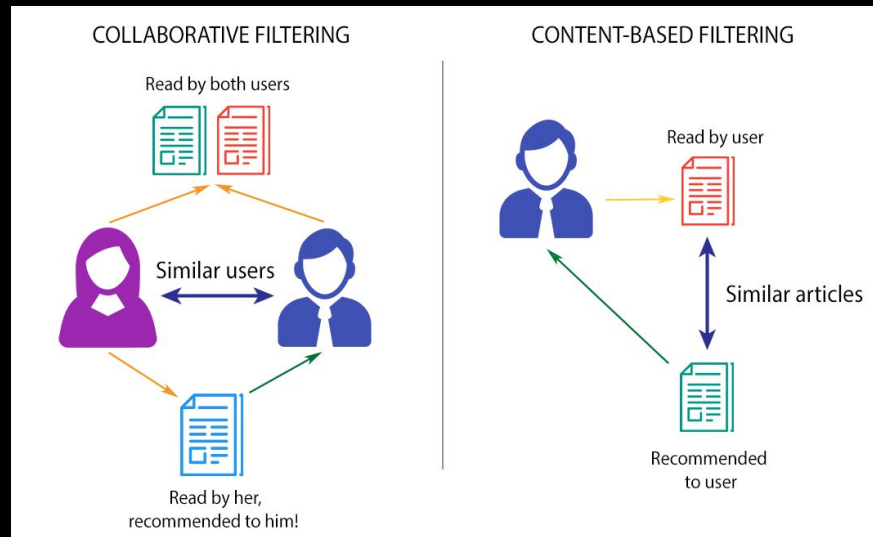
$$y^{(1,1)} = 5$$

$$y^{(1,2)} = 1$$

$$y^{(2,1)} = \text{undefined}$$

Approaches to Recommendations

- **Content-based:** recommends products that are similar products to the ones that the user has liked in the past
- **Collaborative:** recommends products that similar users have liked in the past
- **Hybrid:** a combination of content-based and collaborative



Content-based Filtering

- **Content-based**: recommends products that are similar products to the ones that the user has liked in the past
 - Using the **User - Product** and **Product - Product** relationships
- Types of data used
 - **Product attribute** (necessary)
 - **User behaviour** (at least one “rating”)
- Assume you are given the **feature vector**, we must learn the **user/parameter vector** using linear regression!

Loss & Optimizers

1. Loss



- a. An indicator of how bad the model's prediction was on a single sample

2. Loss/Cost Function

- a. A function which calculates the error in the model across all samples


3. Optimizers

- a. Algorithms which aim to minimize loss by changing the training parameters





Mean Squared Error (Loss Function)

—



$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (y_i - \tilde{y}_i)^2$$

Gradient Descent (optimizer)



repeat until convergence {
 $\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1)$ (for $j = 0$ and $j = 1$)
}

Content-based Filtering Optimization

—

Movie	User 1	User 2	x_1 (comedy)	x_2 (drama)
Step Brothers (1)	5	1	1	0
Anchorman (2)	?	1	1	0
The Other Guys (3)	?	1	1	0
Taxi Driver (4)	4	?	0	1
Goodfellas (5)	5	5	0	1

Dataset with explicit ratings and product attributes

$r(i, j)$: 1 if user j has rated movie i (0 otherwise)

$x^{(i)}$: feature vector for movie i (first index is always 1)

$\theta^{(j)}$: parameter vector for user j

$y^{(i,j)}$: actual rating by user j on movie i (if defined)

$(\theta^{(j)})^T x^{(i)}$: predicted rating by user j on movie i

$$\min_{\theta^{(j)}} \sum_{i:r(i,j)=1} ((\theta^{(j)})^T x^{(i)} - y^{(i,j)})^2$$

Given feature vectors, we are learning the user/parameter vectors

Collaborative Filtering

- **Collaborative**: recommends products that similar users have liked in the past
 - Using the **User - Product** and **User - User** relationships
- Types of data used
 - **User behaviour** (need many “ratings”)
- Assume you are given the **user/parameter vector**, we must learn the **feature vector** using linear regression!

Collaborative Filtering Optimization

Movie	User 1	User 2	x_1 (comedy)	x_2 (drama)
Step Brothers (1)	5	1	?	?
Anchorman (2)	?	1	?	?
The Other Guys (3)	?	1	?	?
Taxi Driver (4)	4	?	?	?
Goodfellas (5)	5	5	?	?

Dataset with explicit ratings and without product attributes

$r(i, j)$: 1 if user j has rated movie i (0 otherwise)

$x^{(i)}$: feature vector for movie i (first index is always 1)

$\theta^{(j)}$: parameter vector for user j

$y^{(i,j)}$: actual rating by user j on movie i (if defined)

$(\theta^{(j)})^T x^{(i)}$: predicted rating by user j on movie i

$$\min_{x^{(i)}} \sum_{i:r(i,j)=1} ((\theta^{(j)})^T x^{(i)} - y^{(i,j)})^2$$

Given user/parameter vectors,
we are learning the feature
vectors

Advantages & Disadvantages

Approach	Advantages	Disadvantages
Content-based	<ul style="list-style-type: none">• Doesn't need many ratings from user before recommendations are usable	<ul style="list-style-type: none">• Different products that haven't been rated by the user or aren't related to other products don't get much exposure• More data is needed
Collaborative	<ul style="list-style-type: none">• Easy to implement• Don't need product attributes/feature vectors	<ul style="list-style-type: none">• User's preferences change, so things they may have liked before may not apply• Needs many ratings from user before recommendations are usable

Movie Recommendation Notebook

Deep Learning Recommender Systems

- Deep learning (DL) refers to machine learning techniques which use artificial neural networks (ANNs)
 - UW DSC Neural Networks Workshop: <https://bit.ly/3B5vgo1>
- DL recommender systems aim to understand the stand relationships discussed and learns the **feature vector** inside the neural networks
- TensorFlow Recommenders Library:
<https://www.tensorflow.org/recommenders>

Resources

- An In-Depth Guide to How Recommender Systems Work:
<https://builtin.com/data-science/recommender-systems>
- Machine Learning by Andrew Ng: <https://www.coursera.org/learn/machine-learning/home/welcome>
- Movie Recommender Systems by Memphis Meng:
<https://towardsdatascience.com/movie-recommendation-system-based-on-movielens-ef0df580cd0e>
- UWaterloo Data Science Channel: https://www.youtube.com/channel/UCknY88pglf2xz_S72WHIDxg